

Projet4_StatsEnGrandeDim

December 12, 2018

1 Projet 4 - Statistiques Grandes Dimensions

```
In [1]: import numpy as np
import pandas as pd
import sklearn.linear_model as skl_linear_mdl
import sklearn.model_selection as skl_mdl_selection
import sklearn.metrics as skl_metrics
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
import random
```

```
In [2]: def print_plot(subtitle,ttitle,xlab,ylab):
    plt.figure()
    plt.suptitle(subtitle,
                 fontsize=14)
    plt.title(ttitle, fontsize=10)
    plt.ylabel(ylab)
    plt.xlabel(xlab)
```

```
In [3]: def print_plot3D(subtitle,ttitle,xlab,ylab,zlab,alpha_rge,lambda_rge,quad_risk,toPivot):
    fig = plt.figure(figsize=(7,7))
    plt.suptitle(subtitle)
    ax = fig.add_subplot(111, projection='3d')
    alpha_axis, lambda_axis = np.meshgrid(alpha_rge, lambda_rge)
    if toPivot==0:
        ax.plot_surface(alpha_axis, lambda_axis, quad_risk.values)
    elif toPivot==1:
        ax.plot_surface(lambda_axis,alpha_axis, all_quad_risk.values)
    ax.set_xlabel(xlab)
    ax.set_ylabel(ylab)
    ax.set_zlabel(zlab)
    plt.show()
```

2 Exercice 1 - Régularisation

Objectif du projet:

Dans ce projet, nous testerons différents estimateurs utilisés dans la prédiction de variables afin d'observer dans quel cas ceux-ci sont les plus appropriés. Les estimateurs sont les suivants: -Lasso -Ridge -Elastic Net

On rappelle: *Lasso*:

$$\hat{\beta}^L = \operatorname{argmin} L_{\lambda}(\beta) \text{ où } L_{\lambda}(\beta) = \frac{1}{2n} \sum_{i=1}^n (Y_i - X_i * \beta) + \lambda ||\beta||_1 \text{ avec } (\lambda > 1)$$

Ridge

$$\hat{\beta}^R = \operatorname{argmin} R_{\mu}(\beta) \text{ où } R_{\mu}(\beta) = \frac{1}{2n} \sum_{i=1}^n (Y_i - X_i * \beta) + \frac{\mu}{2} ||\beta||_2^2 \text{ avec } (\mu > 0)$$

* Elastic Net

$$\beta^{\hat{N}} = \operatorname{argmin} F_{\lambda}(\beta) \text{ où } F_{\lambda}(\beta) = \frac{1}{2n} \sum_{i=1}^n (Y_i - X_i * \beta) + \lambda (\alpha ||\beta||_1 + \frac{1-\alpha}{2} ||\beta||_2^2)$$

3 Question 1

3.0.1 Simulation de X et Y

- On simule une matrice X (1000 * 5000) où chaque coordonnée suit une loi normale centrée réduite
- On simule un vecteur eta de (1000 * 1) où chaque coordonnée suit une loi normale centrée réduite
- On fixe un vecteur beta de (5000 * 1000) où chaque coordonnée est nulle, à l'exception des 15 premières qui valent 1
- On calcul finalement le Y simulé, qui est égale à X * beta + eta
- Enfin, on construit notre dataset, qui est composé du Y en première colonne (c'est la target) et des X sur les 5000 colonnes suivantes (ce sont les features simulées)

```
In [5]: # Simulating X
xindex = ["X" + str(i) for i in range(1, 5001)]
X = pd.DataFrame(np.reshape(np.random.normal(0, 1, (1000*5000)), (1000, 5000)), columns=
# Simulating Eta
eta = pd.DataFrame(np.random.normal(0, 1, 1000), columns=["Eta"])
# Fixing Beta
beta = pd.DataFrame(np.repeat(0, 5000), columns=["Beta"])
for i in range(0, 15):
    beta.loc[i, ] = 1

# Compute simulated Y
Y = pd.DataFrame(np.dot(X, beta) + eta.values, columns=["Y"])
dataset=pd.DataFrame()

dataset = Y
dataset[xindex] = X

dataset.head(5)
```

```

Out [5]:
      Y      X1      X2      X3      X4      X5      X6 \
0 -9.495247  0.252529  0.978173 -1.648376 -0.745150 -0.549010  0.262922
1  1.057552  1.533244  2.006552 -0.190147  0.194154 -0.079179 -2.006603
2  5.028802 -0.166424  2.203721  0.553657 -0.116348  1.172426 -1.116846
3  2.892730 -0.163507  1.075256  0.880031 -0.866662  0.829735 -0.632193
4  9.268786  0.235276  0.227232  0.300307  0.607062  1.144957  0.806097

      X7      X8      X9      ...      X4991      X4992      X4993 \
0 -1.853830 -1.433932 -1.381969  ...      1.249191  0.614602  0.257282
1 -0.960380 -0.295381  0.505217  ...      0.870553  0.931574  1.179348
2  0.268590  0.096138  1.504544  ...      0.253599 -0.174195  1.164225
3  0.224687  2.140223 -0.254601  ...     -1.532971  0.734220 -2.213476
4  1.982821  2.302154 -0.770721  ...      0.209895  0.528178  1.088278

      X4994      X4995      X4996      X4997      X4998      X4999      X5000
0  2.753593  0.191534 -0.458562  0.659595 -1.540402 -1.696091 -0.216240
1 -1.245467  0.303608 -1.480956  2.010872 -0.051878  2.459523 -0.175188
2  1.068989  1.193183 -0.046513  1.621068 -0.521043  0.587093 -0.134657
3 -1.355496 -1.487412 -1.778325  1.236865  0.480267 -0.984885 -0.917793
4  0.738285 -1.541232  1.805350 -0.605179 -0.866061  0.621735 -0.598873

[5 rows x 5001 columns]

```

3.0.2 Séparation du dataset en Train / Test

- On sépare assez classiquement le dataset en 2 parties : le train set, 66% des données le test set, 34% des données

On utilise pour cela la fonction de sklearn : `train_test_split`

```

In [6]: # Splitting X in training and testing set
        train_set, test_set = skl_md1_selection.train_test_split(dataset,
                                                                    train_size=0.66,
                                                                    test_size=0.34)

        X_train = train_set.iloc[:, 1:]
        Y_train = train_set.iloc[:, 0]
        X_test = test_set.iloc[:, 1:]
        Y_test = test_set.iloc[:, 0]

        X_train.head(5)

```

```

Out [6]:
      X1      X2      X3      X4      X5      X6      X7 \
241 -0.411098 -0.145020 -0.601383 -1.016331 -1.211025  0.230308  0.108508
96  -0.909034 -0.605538  0.102225  0.542965  1.159063 -1.391814 -0.927249
518 -1.109685  1.652172 -0.291129  0.912655  0.855644  0.826285 -0.448934
576 -0.837278  1.048874 -1.053501 -0.572636 -1.135570  0.940757  0.164049
35   0.880089 -0.184881  0.117213  0.231647 -1.480395  1.665049  1.736750

```

	X8	X9	X10	...	X4991	X4992	X4993	\
241	1.038159	-0.035076	0.241963	...	-0.216262	0.599253	-1.062285	
96	-0.325360	-1.691028	-1.453019	...	1.991209	-0.032368	-1.259970	
518	-1.040584	-1.312625	-0.082198	...	0.820791	2.303144	0.472122	
576	-1.177026	0.284111	-0.776715	...	-0.873612	-0.215861	-0.339654	
35	-0.729738	1.122556	-1.219183	...	0.592456	1.162374	0.742223	

	X4994	X4995	X4996	X4997	X4998	X4999	X5000
241	-0.844311	-2.037630	1.039755	0.369866	0.073692	2.185047	-0.961506
96	0.302735	1.830717	-0.267040	1.010104	1.145018	1.549268	1.234028
518	0.361112	-0.138547	-0.385430	-0.720588	1.084893	-0.397409	-0.706167
576	1.017378	-0.062421	-2.025236	-1.283571	-1.305765	0.573330	-0.256140
35	-0.217050	-1.052846	0.237344	-1.439791	-0.915768	-0.643530	0.887570

[5 rows x 5000 columns]

3.0.3 A) On estime la Régression Elastic Net pour alpha variant de 0 à 1

On estime les 5000 coefficients pour alpha variant entre 0 et 1 (lambda est égal à 1): * on crée un range de alpha de taille 11 * on itère sur ce range, en calculant à chaque itération le vecteur beta estimé de (5000) on remplit un dictionnaire qui contiendra les 5000 * 11 betas issus de nos simulations

```
In [7]: def elastic_net_return_coeff(iterateOnL1 ,alpha_value, l1ratio_value, X_train, Y_train):
    all_beta_chap = dict()
    toIterate = np.linspace(0, 1, 11)
    for i_toIterate in toIterate:

        if iterateOnL1==True:
            lasso = skl_linear_md1.ElasticNet(alpha=alpha_value, l1_ratio=i_toIterate,
                                                fit_intercept=False)#, max_iter = 10000)

        if iterateOnL1==False:
            lasso = skl_linear_md1.ElasticNet(alpha=i_toIterate, l1_ratio=l1ratio_value,
                                                fit_intercept=False)#, max_iter = 10000)

        lasso.fit(X_train, Y_train)
        beta_chap = lasso.coef_

        all_beta_chap.update({i_toIterate: beta_chap})

    return all_beta_chap
```

- Affichage des coefficients pour chaque valeur de alpha

On affiche, pour chacun des 5000 betas, sa valeur en fonction de alpha On peut donc observer dans quelle mesure le modèle shrinke les coefficients en fonction de la valeur de alpha. Dans notre cas, on sais que 15 coefficients sont égaux à 1, tandis que les 4985 autres sont nuls : * on observe sur le graphique que pour un alpha entre 0 et 0.2, nos 15 coefficients valent à peu près 0.5, et que

leur variance diminue à mesure que alpha augmente. Cependant, ils sont en moyenne de plus en plus éloigné de leurs vrais valeurs (à savoir, 1), et sont tous nuls lorsque alpha vaut 1. * En ce qui concerne les 4985 coefficients nuls, leur variance est élevée pour un alpha proche de zéro (et ils ne sont donc pas tous nuls). En revanche, dès que alpha augmente, ils deviennent tous nuls (leurs vraie valeur)

La bonne valeur de alpha est donc un peu éloigné de 0 (afin que les 4985 coefficients nuls

```
In [8]: all_beta_chap = dict()

all_beta_chap = elastic_net_return_coeff(iterateOnL1=True,alpha_value=1,l1ratio_value=
X_train=X_train,Y_train=Y_train)

beta_vs_alpha = pd.DataFrame(all_beta_chap)
print_plot("Regularization Path for Elastic Net over Alpha Parameter","Lambda = 1","Al
alpha_range=np.linspace(0,1,11)
for row in beta_vs_alpha.iterrows():
    row = row[1]
    plt.plot(alpha_range, row)

beta_vs_alpha.head(20)


/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_desc
ConvergenceWarning)
```

```
Out [8]:
```

	0.0	0.1	0.2	0.3	0.4	0.5	0.6	\
0	0.100806	0.323203	0.417339	0.409784	0.373449	0.331859	0.284457	
1	0.106871	0.327132	0.413363	0.406178	0.370956	0.329974	0.283242	
2	0.074613	0.237693	0.306034	0.284478	0.230360	0.166249	0.090825	
3	0.100188	0.283142	0.341606	0.319563	0.272003	0.216055	0.150631	
4	0.120269	0.324353	0.421593	0.413410	0.382287	0.347236	0.308203	
5	0.104607	0.301891	0.391174	0.379363	0.339577	0.293497	0.240428	
6	0.119875	0.351485	0.442831	0.432993	0.399694	0.361618	0.318276	
7	0.123641	0.363620	0.458348	0.455737	0.425183	0.389203	0.348364	
8	0.149943	0.418535	0.499025	0.489648	0.463071	0.433925	0.401651	
9	0.101699	0.298290	0.379534	0.360586	0.317404	0.267062	0.208529	
10	0.114183	0.316986	0.414996	0.408583	0.370037	0.325454	0.274041	
11	0.093782	0.281892	0.361690	0.339886	0.294738	0.242542	0.181977	
12	0.111085	0.321473	0.406366	0.393105	0.355100	0.311764	0.262063	
13	0.102499	0.304409	0.380443	0.366291	0.325590	0.277355	0.221610	
14	0.121217	0.354160	0.446415	0.440141	0.407051	0.368502	0.324577	
15	-0.013283	-0.000000	-0.000000	-0.000000	-0.000000	-0.000000	-0.000000	
16	0.020846	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
17	0.013522	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
18	-0.000957	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
19	0.032096	0.029276	0.000000	0.000000	0.000000	0.000000	0.000000	
	0.7	0.8	0.9	1.0				

0	0.229974	0.165231	0.086854	0.000000
1	0.229451	0.170432	0.100176	0.013836
2	0.000873	0.000000	0.000000	0.000000
3	0.073064	0.000000	0.000000	0.000000
4	0.264556	0.214905	0.152503	0.076177
5	0.178629	0.105620	0.020721	0.000000
6	0.268459	0.213713	0.144600	0.070305
7	0.301591	0.250769	0.196183	0.126612
8	0.365799	0.328105	0.281959	0.218631
9	0.139578	0.065072	0.000000	0.000000
10	0.214092	0.144151	0.057395	0.000000
11	0.110845	0.027127	0.000000	0.000000
12	0.204448	0.143290	0.069195	0.000000
13	0.156431	0.080252	0.000000	0.000000
14	0.274067	0.211117	0.140981	0.061083
15	-0.000000	-0.000000	-0.000000	-0.000000
16	0.000000	0.000000	0.000000	0.000000
17	0.000000	0.000000	0.000000	0.000000
18	0.000000	0.000000	-0.000000	-0.000000
19	0.000000	0.000000	0.000000	0.000000

Regularization Path for Elastic Net over Alpha Parameter
Lambda = 1



3.0.4 B) Chemin de Régularisation (Regularization Path) de l'estimateur Lasso

On estime les 5000 coefficients pour lambda variant entre 0 et 1 (alpha est égal à 1): * on crée un range de lambda de taille 11 * on itère sur ce range, en calculant à chaque itération le vecteur beta estimé de (50001) on remplit un dictionnaire qui contiendra les 5000 * 11 betas issus de nos simulations

De la même manière que pour le coefficient alpha, on affiche, pour chacun des 5000 betas, sa valeur en fonction de lambda : * On voit qu'entre 0 et 0.2, les 15 coefficients sont très proches de leurs vrai valeurs (1), puis descendent linéairement jusqu'à 0, lorsque lambda vaut 1 * les 4985 nuls sont un peu dispersés autour de 0, pour alpha entre 0 et 0.2. Après 0.2, les 4985 coefficients estimés sont tous nuls (égaux à leurs vrai valeur, donc)

Graphiquement, on poserais un lambda égal à 0.1 ou 0.2

- Méthode 1 : En calculant nous mêmes

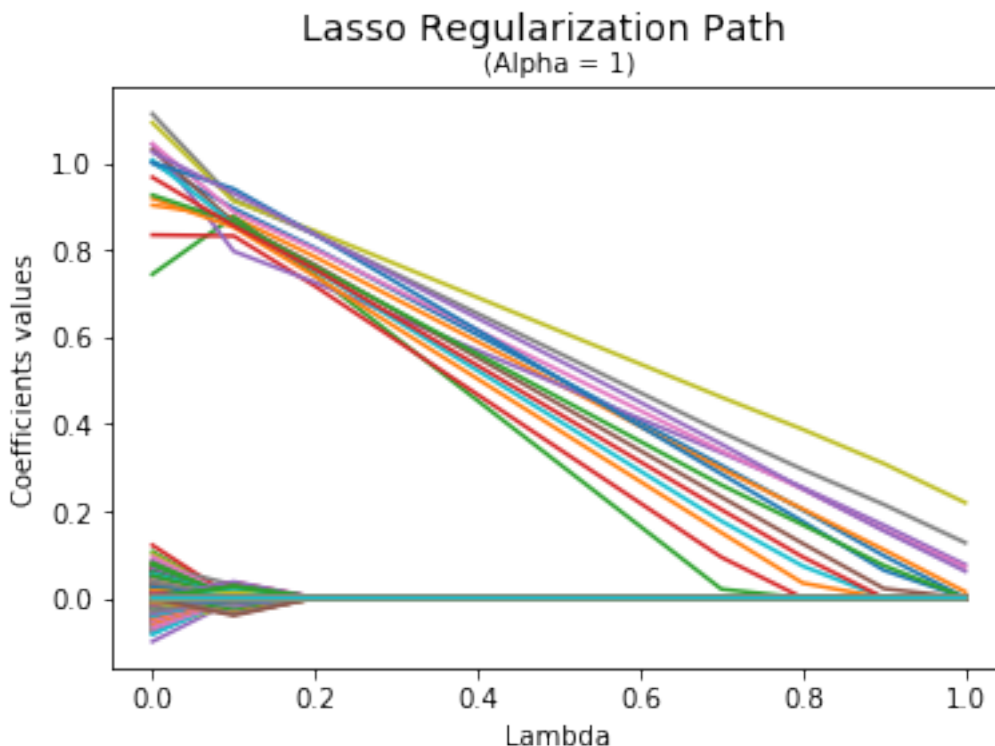
```
In [141]: def print_regularization_path(X_train,Y_train,isLasso=False):
    all_quad_risk = dict()
    all_beta_chap = dict()
    lambda_range = np.linspace(0, 1, 11)
    if isLasso==True:
        all_beta_chap = elastic_net_return_coeff(False,1,1,X_train,Y_train)
    if isLasso==False:
        all_beta_chap = elastic_net_return_coeff(False,1,0,X_train,Y_train)

    beta_vs_lambda = pd.DataFrame(all_beta_chap)
    if isLasso==True:
        print_plot("Lasso Regularization Path","(Alpha = 1)","Lambda","Coefficients v
    if isLasso==False:
        print_plot("Ridge Regularization Path","(Alpha = 0)","Lambda","Coefficients v
    for row in beta_vs_lambda.iterrows():
        row = row[1]
        plt.plot(lambda_range, row)
```

```
In [10]: print_regularization_path(X_train,Y_train,isLasso=True)
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:14: UserWarning
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_desc
positive)
```



- Méthode 2 : Avec la méthode Path

```
In [11]: #lasso = skl_linear_mdl.ElasticNet(l1_ratio=1)
#lasso_path = lasso.path(X_train, Y_train, l1_ratio=1)
#lasso_path_result = lasso_path[1]
#print_plot("", "Lasso Regularization Path", "Lambda", "Coefficients values")

#for i in range(0, len(lasso_path_result)):
#    plt.plot(lasso_path[0], lasso_path_result[i])
```

3.0.5 B.1) Chemin de Régularisation de l'estimateur Ridge


De la même manière que pour l'estimateur Lasso, on affiche, pour chacun des 5000 coefficients du modèle Ridge, leur valeur en fonction de Lambda (entre 0 et 1), pour un alpha égal à 0 (par définition du modèle Ridge) On peut observer que les 15 coefficients unitaires tendent beaucoup plus vite vers 0 (à 0.1, ils sont déjà presque nuls)

```
In [12]: print_regularization_path(X_train, Y_train, False)
```

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:14: UserWarning

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:141: UserWarning: Warning: The 'alpha' parameter is deprecated in favor of 'lambda_1' and 'lambda_2'. It will be removed in a future version of scikit-learn.


```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:455: ConvergenceWarning
```



3.0.6 C) Optimisation du / des paramètres de régularisation (lambda et alpha)

Nous allons maintenant, pour chacun des 3 modèles, optimiser leurs paramètres de régularisation :

- * Pour le Lasso, on optimise donc pour lambda entre 0 et 1, sachant qu'alpha vaut 1 (par définition du modèle Lasso)
- * Pour le Ridge, on optimise donc pour lambda entre 0 et 1, sachant qu'alpha vaut 0 (par définition du Ridge)
- * Pour l'Elastic Net, qui est une pondération des modèles Lasso et Ridge, on optimise donc sur lambda et sur alpha, variant tout deux entre 0 et 1

Nous optimiserons ces paramètres par validation croisée :

- * estimation du modèle sur le training set
- * calcul du MSE par rapport au test set

```
In [13]: def elastic_net_return_error(X_train, Y_train, X_test, Y_test, alpha_range, lambda_range):
    all_quad_risk = pd.DataFrame(index=alpha_range,
                                columns=lambda_range)

    for i_alpha in alpha_range:
        for i_lambda in lambda_range:
            en = skl_linear_md1.ElasticNet(alpha=i_lambda, l1_ratio=i_alpha,
                                           fit_intercept = False, max_iter = 10000)
            en.fit(X_train, Y_train)
```

```

        beta_chap = en.coef_
        Y_predict = en.predict(X_test)
        quad_risk = skl_metrics.mean_squared_error(Y_test, Y_predict)
        all_quad_risk.loc[i_alpha, i_lambda] = quad_risk

min_mse = all_quad_risk.min().min()
x, y = np.where(all_quad_risk.values == min_mse)

alpha_opt_elastic_net = all_quad_risk.index[x].tolist()[0]
#lambda_opt_elastic_net = all_quad_risk.index[x].tolist()[0]
lambda_opt_elastic_net = all_quad_risk.columns[y].tolist()[0]

return lambda_opt_elastic_net, alpha_opt_elastic_net, all_quad_risk

In [14]: def optimise_params(X_train, Y_train, X_test, Y_test, model_num):
    all_quad_risk = dict()

    if model_num == 1:
        alpha_opt = 1
        #lambda_opt, all_quad_risk=Elastic_Net_return_error(1, X_train, Y_train, X_test)
        lambda_opt, _, all_quad_risk = elastic_net_return_error(X_train, Y_train, X_test,
                                                                [alpha_opt], np.linspace(0, 1, 100))

    elif model_num == 2:
        alpha_opt = 0
        #lambda_opt, all_quad_risk=Elastic_Net_return_error(0, X_train, Y_train, X_test)
        lambda_opt, _, all_quad_risk = elastic_net_return_error(X_train, Y_train, X_test,
                                                                [alpha_opt], np.linspace(0, 1, 100))

    elif model_num == 3:
        lambda_opt, alpha_opt, all_quad_risk = elastic_net_return_error(X_train, Y_train, X_test,
                                                                np.linspace(0, 1, 100))

    return lambda_opt, alpha_opt, all_quad_risk

def print_opt_model(alpha_opt, lambda_opt, model_num):
    if model_num == 1:
        print_plot("Lasso Regression", "Lambda Optimal : " + str(lambda_opt), "Lambda", "MSE")
    elif model_num == 2:
        print_plot("Ridge Regression", "Lambda Optimal : " + str(lambda_opt), "Lambda", "MSE")
    elif model_num == 3:
        pass

```

- Lasso

```


In [15]: lambda_opt_lasso, _, all_quad_risk = optimise_params(X_train, Y_train, X_test, Y_test)
        print_plot("Lasso Regression", "Lambda Optimal : " + str(lambda_opt_lasso), "Lambda", "MSE")
        plt.plot(all_quad_risk.columns,
                 all_quad_risk.values.reshape(11, 1))

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning
if __name__ == '__main__':

```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:107: UserWarning:
  positive)
```

```
Out[15]: [ <matplotlib.lines.Line2D at 0x1a14aaacc0>]
```




- Ridge

```
In [16]: lambda_opt_ridge, _, all_quad_risk = optimise_params(X_train, Y_train, X_test, Y_test)
print_plot("Ridge Regression", "Lambda Optimal : " + str(lambda_opt_ridge), "Lambda", "MSE")
plt.plot(all_quad_risk.columns,
         all_quad_risk.values.reshape(11, 1))
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning:
  if __name__ == '__main__':
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:107: UserWarning:
  positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:107: ConvergenceWarning:
```

```
Out[16]: [ <matplotlib.lines.Line2D at 0x1a14732978>]
```



- Elastic Net

```
In [17]: lambda_opt_elastic_net, alpha_opt_elastic_net, all_quad_risk = optimise_params(X_train


print_plot3D("Elastic Net Regression\nAlpha Optimal : " + str(alpha_opt_elastic_net) +
"\nLambda Optimal : " + str(lambda_opt_elastic_net), "", "Alpha", "Lamda", "Risk",
np.linspace(0, 1, 11), np.linspace(0, 1, 11), all_quad_risk, 0)

print_plot3D("Elastic Net Regression\nAlpha Optimal : " + str(alpha_opt_elastic_net) +
"\nLambda Optimal : " + str(lambda_opt_elastic_net), "", "Lambda", "Alpha", "Risk",
np.linspace(0, 1, 11), np.linspace(0, 1, 11), all_quad_risk, 1)


opt_elastic_net=[]
opt_elastic_net.append(alpha_opt_elastic_net)
opt_elastic_net.append(lambda_opt_elastic_net)

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning
if __name__ == '__main__':
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100:
positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100:
ConvergenceWarning)
```

Elastic Net Regression
Alpha Optimal : 1.0
Lambda Optimal : 0.1



Elastic Net Regression
Alpha Optimal : 1.0
Lambda Optimal : 0.1



3.0.7 D) Score des estimateurs

Enfin, on sélectionne le meilleurs des 3 modèles, en sélectionnant celui qui affiche le MSE le plus faible. On prend également soin de calculer le coefficient de détermination R2 du modèle.

```
In [18]: def score(X_train,Y_train, lambda_opt,isLasso):  
    if isLasso==True:  
        model = skl_linear_md1.ElasticNet(alpha=lambda_opt, l1_ratio=1,  
                                           fit_intercept=False, max_iter = 10000)  
    if isLasso==False:  
        model = skl_linear_md1.ElasticNet(alpha=lambda_opt, l1_ratio=0,
```

```

fit_intercept=False, max_iter = 10000)
model.fit(X_train, Y_train)
score_model = model.score(X_test, Y_test)
Y_predict = model.predict(X_test)
mse_model = skl_metrics.mean_squared_error(Y_test, Y_predict)

return[score_model,mse_model]

In [19]: def score_elastic_net(X_train,Y_train, opt_elastic_net):
en = skl_linear_mdl.ElasticNet(alpha=opt_elastic_net[1], l1_ratio=opt_elastic_net
fit_intercept=False, max_iter = 10000)

en.fit(X_train, Y_train)
score_en = en.score(X_test, Y_test)
Y_predict_en = en.predict(X_test)
mse_en = skl_metrics.mean_squared_error(Y_test, Y_predict_en)

return([score_en,mse_en])

In [20]: def print_score(X_train, Y_train, lambda_opt_lasso, lambda_opt_ridge, opt_elastic_net):
lasso=score(X_train,Y_train,lambda_opt=lambda_opt_lasso,isLasso=True)
score_lasso=lasso[0]
mse_lasso=lasso[1]
ridge=score(X_train,Y_train,lambda_opt=lambda_opt_ridge,isLasso=False)
score_ridge=ridge[0]
mse_ridge=ridge[1]
en=score_elastic_net(X_train,Y_train,opt_elastic_net=opt_elastic_net)
score_en=en[0]
mse_en=en[1]

df_score = pd.DataFrame([score_lasso, score_ridge, score_en],
columns=["Score (R2)"], index=["Lasso", "Ridge", "Elastic Net"])
df_score["MSE"] = [mse_lasso, mse_ridge, mse_en]

best_model_R2 = df_score.index[np.argmax(df_score["Score (R2)"].values)]
best_model_mse = df_score.index[np.argmin(df_score["MSE"].values)]
df_score.loc["BEST MODEL", "Score (R2)"] = best_model_R2
df_score.loc["BEST MODEL", "MSE"] = best_model_mse
print(df_score.head())

In [21]: print_score(X_train,Y_train,lambda_opt_lasso,lambda_opt_ridge,opt_elastic_net)

Score (R2)      MSE
Lasso      0.912584  1.19948
Ridge      0.821044  2.45554
Elastic Net 0.912584  1.19948
BEST MODEL      Lasso      Lasso

```

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:8: UserWarning

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:107: DeprecationWarning:
positive)
```

4 Question 2

On reproduit les étapes de la Question 1, sur un Dataset différent. Pour chacun des 3 modèles : *

On affiche le regularization Path * On optimise le/les paramètre(s) de régularisation

Enfin, on sélectionne le meilleur modèle

```
In [22]: # Simulating X
xindex = ["X" + str(i) for i in range(1, 5001)]
X = pd.DataFrame(np.reshape(np.random.normal(0, 1, (1000*5000)), (1000, 5000)),
                  columns=xindex)

# Simulating Eta
eta = pd.DataFrame(np.random.normal(0, 1, 1000),
                   columns=["Eta"])

# Simulating Beta
beta = pd.DataFrame(np.repeat(0, 5000),
                    columns=["Beta"])
for i in range(0, 1500):
    beta.loc[i, ] = 1

# Compute simulated Y
Y = pd.DataFrame(np.dot(X, beta) + eta.values,
                  columns=["Y"])

dataset=pd.DataFrame()
dataset = Y
dataset[xindex] = X

# Splitting X in training and testing set
train_set, test_set = skl_md1_selection.train_test_split(dataset,
                                                           train_size=0.66,
                                                           test_size=0.34)

X_train = train_set.iloc[:, 1:]
Y_train = train_set.iloc[:, 0]
X_test = test_set.iloc[:, 1:]
Y_test = test_set.iloc[:, 0]
```


- Lasso


```
In [23]: print_regularization_path(X_train,Y_train,isLasso=True)
lambda_opt_lasso, _, all_quad_risk = optimise_params(X_train, Y_train, X_test, Y_test)
print_plot("Lasso Regression","Lambda Optimal : " + str(lambda_opt_lasso),"Lambda","M
plt.plot(all_quad_risk.columns,
         all_quad_risk.values.reshape(11, 1))
```



```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:14: UserWarning  
  
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:141: UserWarning:   
    positive)  
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning:   
    if __name__ == '__main__':
```

Out[23]: [





- Ridge

```
In [24]: print_regularization_path(X_train,Y_train,isLasso=False)
lambda_opt_ridge, _, all_quad_risk = optimise_params(X_train, Y_train, X_test, Y_test)
print_plot("Ridge Regression","Lambda Optimal : " + str(lambda_opt_ridge),"Lambda","MSE")
plt.plot(all_quad_risk.columns,
         all_quad_risk.values.reshape(11, 1))
```

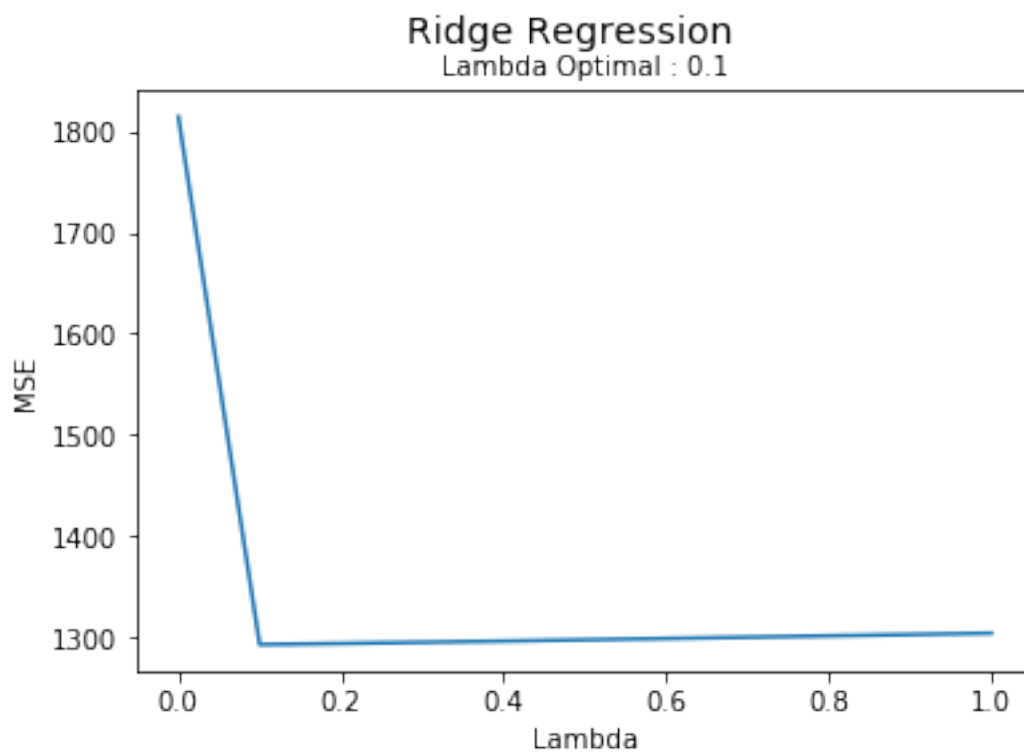
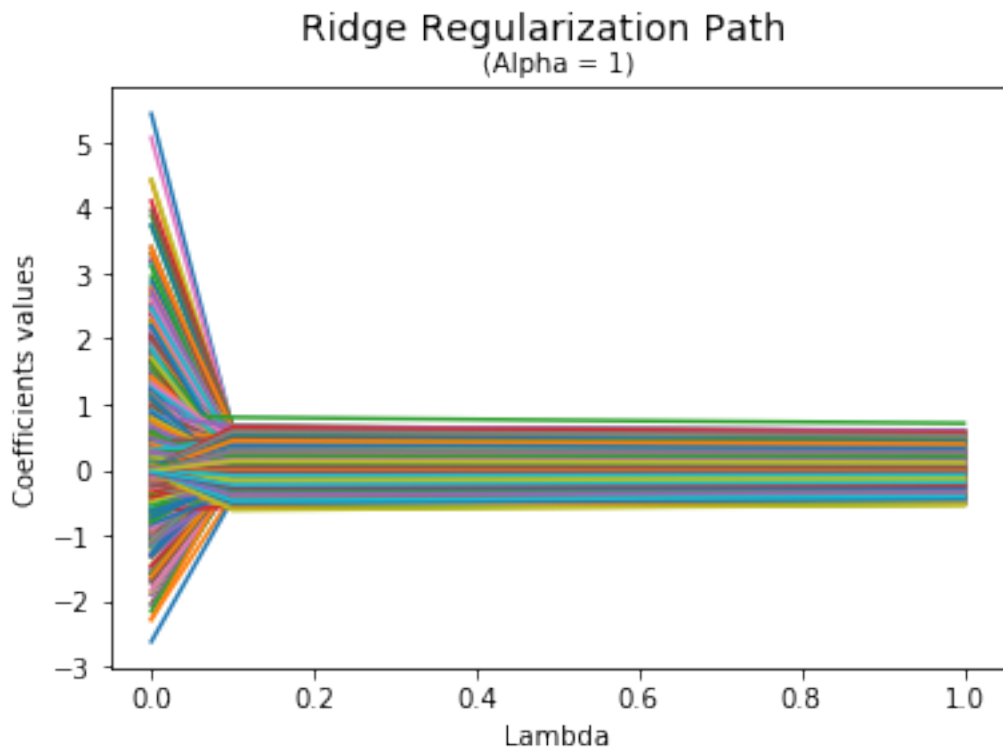
```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:14: UserWarning
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100:
positive)
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_desc
ConvergenceWarning)
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning
  if __name__ == '__main__':
```

```
Out[24]: []
```



- Elastic Net

```
In [25]: lambda_opt_elastic_net, alpha_opt_elastic_net, all_quad_risk = optimise_params(X_train, y_train)


print_plot3D("Elastic Net Regression\nAlpha Optimal : " + str(alpha_opt_elastic_net) + "\nLambda Optimal : " + str(lambda_opt_elastic_net), "", "Alpha", "Lambda", "Risk",
            np.linspace(0, 1, 11), np.linspace(0, 1, 11), all_quad_risk, 0)

print_plot3D("Elastic Net Regression\nAlpha Optimal : " + str(alpha_opt_elastic_net) + "\nLambda Optimal : " + str(lambda_opt_elastic_net), "", "Lambda", "Alpha", "Risk",
            np.linspace(0, 1, 11), np.linspace(0, 1, 11), all_quad_risk, 1)


opt_elastic_net=[]
opt_elastic_net.append(alpha_opt_elastic_net)
opt_elastic_net.append(lambda_opt_elastic_net)

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning:
    if __name__ == '__main__':
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100: ConvergenceWarning:
    positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100: ConvergenceWarning)
```

Elastic Net Regression
Alpha Optimal : 0.0
Lambda Optimal : 0.1



Elastic Net Regression
 Alpha Optimal : 0.0
 Lambda Optimal : 0.1



- Score

```
In [26]: print_score(X_train,Y_train,lambda_opt_lasso,lambda_opt_ridge,opt_elastic_net)
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:455: ConvergenceWarning
```

	Score (R2)	MSE
Lasso	-0.069466	1647.65
Ridge	0.160783	1292.92

Elastic Net	0.160783	1292.92
BEST MODEL	Ridge	Ridge

5 Question 3

On reproduit les étapes de la Question 1, sur un Dataset différent. Pour chacun des 3 modèles : *

On affiche le regularization Path * On optimise le/les paramètre(s) de régularisation

Enfin, on sélectionne le meilleur modèle

```
In [33]: # Simulating X
xindex = ["X" + str(i) for i in range(1, 51)]
cov_mat = np.repeat(0.7, 50*50).reshape(50, 50)
for i in range(0, 50):
    for j in range(i, 50):
        cov_mat[i, j] = np.power(cov_mat[i, j], abs(i-j))
        cov_mat[j, i] = np.power(cov_mat[j, i], abs(i-j))

X = pd.DataFrame(np.reshape(np.random.multivariate_normal([0] * 50, cov_mat, 100),
                           (100, 50)),
                 columns=xindex)

# Simulating Eta
eta = pd.DataFrame(np.random.normal(0, 1, 100),
                  columns=["Eta"])

# Simulating Beta
beta = pd.DataFrame(np.repeat(0, 50),
                   columns=["Beta"])

beta.loc[0, ] = 10
beta.loc[1, ] = 10
beta.loc[2, ] = 5
beta.loc[3, ] = 5
for i in range(4, 14):
    beta.loc[i, ] = 1

# Compute simulated Y
Y = pd.DataFrame(np.dot(X, beta) + eta.values,
                 columns=["Y"])

dataset = pd.DataFrame()

dataset = Y
dataset[xindex] = X

# Splitting X in training and testing set
```

```
train_set, test_set = skl_mdl_selection.train_test_split(dataset,
                                                         train_size=0.66,
                                                         test_size=0.34)
```

```
X_train = train_set.iloc[:, 1:]
Y_train = train_set.iloc[:, 0]
X_test = test_set.iloc[:, 1:]
Y_test = test_set.iloc[:, 0]
```

- Lasso

```
In [34]: print_regularization_path(X_train,Y_train,isLasso=True)
        lambda_opt_lasso, _, all_quad_risk = optimise_params(X_train, Y_train, X_test, Y_test)
        print_plot("Lasso Regression","Lambda Optimal : " + str(lambda_opt_lasso),"Lambda","MSE")
        plt.plot(all_quad_risk.columns,
                 all_quad_risk.values.reshape(11, 1))
```


/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:14: UserWarning


/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:141: UserWarning:
 positive)

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:141: UserWarning:
 ConvergenceWarning)

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning
 if __name__ == '__main__':

Out[34]: [<matplotlib.lines.Line2D at 0x1a1393eeb8>]







- Ridge

```
In [35]: print_regularization_path(X_train,Y_train,isLasso=False)
         lambda_opt_ridge, _, all_quad_risk = optimise_params(X_train, Y_train, X_test, Y_test)
         print_plot("Ridge Regression","Lambda Optimal : " + str(lambda_opt_ridge),"Lambda","MSE")
         plt.plot(all_quad_risk.columns,
                  all_quad_risk.values.reshape(11, 1))

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:14: UserWarning:
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent:
positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent:
ConvergenceWarning)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning:
if __name__ == '__main__':
```

```
Out[35]: [matplotlib.lines.Line2D at 0x1a2049c780>]
```



- Elastic Net

```
In [36]: lambda_opt_elastic_net, alpha_opt_elastic_net, all_quad_risk = optimise_params(X_train, y_train)


print_plot3D("Elastic Net Regression\nAlpha Optimal : " + str(alpha_opt_elastic_net) + "\nLambda Optimal : " + str(lambda_opt_elastic_net), "", "Alpha", "Lambda", "Risk",
            np.linspace(0, 1, 11), np.linspace(0, 1, 11), all_quad_risk, 0)

print_plot3D("Elastic Net Regression\nAlpha Optimal : " + str(alpha_opt_elastic_net) + "\nLambda Optimal : " + str(lambda_opt_elastic_net), "", "Lambda", "Alpha", "Risk",
            np.linspace(0, 1, 11), np.linspace(0, 1, 11), all_quad_risk, 1)


opt_elastic_net=[]
opt_elastic_net.append(alpha_opt_elastic_net)
opt_elastic_net.append(lambda_opt_elastic_net)

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning:
    if __name__ == '__main__':
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100: ConvergenceWarning:
    positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100: ConvergenceWarning)
```

Elastic Net Regression
Alpha Optimal : 0.8
Lambda Optimal : 0.1



Elastic Net Regression
Alpha Optimal : 0.8
Lambda Optimal : 0.1



- Score

```
In [37]: print_score(X_train,Y_train,lambda_opt_lasso,lambda_opt_ridge,opt_elastic_net)
```

	Score (R2)	MSE
Lasso	0.997821	1.68367
Ridge	0.993142	5.30039
Elastic Net	0.998388	1.24564
BEST MODEL	Elastic Net	Elastic Net

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:8: UserWarning

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100: ConvergenceWarning:
  positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100: ConvergenceWarning)
```

6 Exercice 2

Pour l'application à des données réelles, on choisi de partir d'un dataset sur le Bitcoin d'une quinzaine de features, puis de faire du feature engineering sur ces features afin d'avoir une matrice de features de taille très importante (comme pour l'exercice 1).

On décide également d'introduire des features qui, vraisemblablement, n'auront pas d'effet sur la target : leurs coefficients de régression seront à priori nuls ! (et c'est ce qu'on souhaite).

On se retrouve finalement avec une matrice de (1337 * 127), qui parait être un bon candidat pour une régression pénalisée.

6.1 Initialisation

- Import des données

```
In [123]: import datetime
```

```
df = pd.read_csv("bitcoin_dataset.csv")
df_2 = pd.read_csv("bitcoin_price.csv")
df_3 = pd.read_csv("litecoin_price.csv")
df_4 = pd.read_csv("monero_price.csv")

df_2 = df_2[:::-1]
df_2 = df_2.reset_index(drop=True)

df_3 = df_3[:::-1]
df_3 = df_3.reset_index(drop=True)

df_4 = df_4[:::-1]
df_4 = df_4.reset_index(drop=True)

df["Date"] = df["Date"].apply(lambda x: datetime.datetime.strptime(x, "%Y-%m-%d %H:%M:%S"))
df_2["Date"] = df_2["Date"].apply(lambda x: datetime.datetime.strptime(x, "%b %d, %Y"))
df_3["Date"] = df_3["Date"].apply(lambda x: datetime.datetime.strptime(x, "%b %d, %Y"))
df_4["Date"] = df_4["Date"].apply(lambda x: datetime.datetime.strptime(x, "%b %d, %Y"))

df_3.columns = ["Date", "Open_Litecoin", "High_Litecoin", "Low_Litecoin",
               "Close_Litecoin", "Volume_Litecoin", "Market Cap_Litecoin"]

df_4.columns = ["Date", "Open_Monero", "High_Monero", "Low_Monero",
               "Close_Monero", "Volume_Monero", "Market Cap_Monero"]

df = df.merge(df_2)
```

```
df = df.merge(df_3)
df = df.merge(df_4)
df.head(5)
```

```
Out[123]:
```

	Date	btc_market_price	btc_total_bitcoins	btc_market_cap	\
0	2014-05-21	494.87	12798675.0	6.333680e+09	
1	2014-05-22	523.84	12802850.0	6.706645e+09	
2	2014-05-23	527.47	12806800.0	6.755203e+09	
3	2014-05-24	521.52	12811000.0	6.681193e+09	
4	2014-05-25	575.00	12814775.0	7.368496e+09	

	btc_trade_volume	btc_blocks_size	btc_avg_block_size	\
0	1.185312e+07	18419.0	0.250119	
1	2.358140e+07	18459.0	0.241909	
2	2.826269e+07	18501.0	0.271977	
3	9.312733e+06	18533.0	0.195074	
4	3.730564e+07	18566.0	0.219971	

	btc_n_orphaned_blocks	btc_n_transactions_per_block	\
0	0.0	389.0	
1	2.0	519.0	
2	0.0	408.0	
3	0.0	503.0	
4	1.0	410.0	

	btc_median_confirmation_time	...	Low_Litecoin	\
0	6.533333	...	10.53	
1	7.450000	...	10.48	
2	6.933333	...	10.65	
3	7.350000	...	10.99	
4	8.466667	...	11.12	

	Close_Litecoin	Volume_Litecoin	Market	Cap_Litecoin	Open_Monero	\
0	10.59	2,653,860		302,609,000	2.47	
1	10.67	5,186,880		302,201,000	1.59	
2	11.23	9,134,640		304,423,000	2.05	
3	11.10	2,574,380		320,801,000	2.92	
4	11.53	9,072,890		318,703,000	4.04	

	High_Monero	Low_Monero	Close_Monero	Volume_Monero	Market	Cap_Monero
0	2.65	1.23	1.60	246,540		2,079,640
1	2.19	1.36	2.10	132,918		1,371,470
2	3.43	2.05	2.96	266,852		1,816,200
3	4.01	2.62	3.70	248,028		2,653,720
4	4.04	2.80	3.14	283,545		3,774,890

[5 rows x 42 columns]

- Feature Engineering

```

In [124]: df["Open_Return"] = df["Open"].pct_change()
df["High_Return"] = df["High"].pct_change()
df["Low_Return"] = df["Low"].pct_change()
df["Close_Return"] = df["Close"].pct_change()

df["Open_Return_Litecoin"] = df["Open_Litecoin"].pct_change()
df["High_Return_Litecoin"] = df["High_Litecoin"].pct_change()
df["Low_Return_Litecoin"] = df["Low_Litecoin"].pct_change()
df["Close_Return_Litecoin"] = df["Close_Litecoin"].pct_change()

df["Open_Return_Monero"] = df["Open_Monero"].pct_change()
df["High_Return_Monero"] = df["High_Monero"].pct_change()
df["Low_Return_Monero"] = df["Low_Monero"].pct_change()
df["Close_Return_Monero"] = df["Close_Monero"].pct_change()

df["Volume"] = df["Volume"].apply(lambda x: float(x.replace(",", "")) if x != "-" else 0)
df["Market Cap"] = df["Market Cap"].apply(lambda x: float(x.replace(",", "")) if x != "-" else 0)

df["Volume_Litecoin"] = df["Volume_Litecoin"].apply(lambda x: float(x.replace(",", "")) if x != "-" else 0)
df["Market Cap_Litecoin"] = df["Market Cap_Litecoin"].apply(lambda x: float(x.replace(",", "")) if x != "-" else 0)

df["Volume_Monero"] = df["Volume_Monero"].apply(lambda x: float(x.replace(",", "")) if x != "-" else 0)
df["Market Cap_Monero"] = df["Market Cap_Monero"].apply(lambda x: float(x.replace(",", "")) if x != "-" else 0)

# Moyenne Mobile
for i_lag in [3, 6, 9]:
    df["Open_Return_MA" + str(i_lag) + "D"] = df["Open_Return"].rolling(i_lag).mean()
    df["High_Return_MA" + str(i_lag) + "D"] = df["High_Return"].rolling(i_lag).mean()
    df["Low_Return_MA" + str(i_lag) + "D"] = df["Low_Return"].rolling(i_lag).mean()
    df["Close_Return_MA" + str(i_lag) + "D"] = df["Close_Return"].rolling(i_lag).mean()

    df["Open_Return_Litecoin_MA" + str(i_lag) + "D"] = df["Open_Return_Litecoin"].rolling(i_lag).mean()
    df["High_Return_Litecoin_MA" + str(i_lag) + "D"] = df["High_Return_Litecoin"].rolling(i_lag).mean()
    df["Low_Return_Litecoin_MA" + str(i_lag) + "D"] = df["Low_Return_Litecoin"].rolling(i_lag).mean()
    df["Close_Return_Litecoin_MA" + str(i_lag) + "D"] = df["Close_Return_Litecoin"].rolling(i_lag).mean()

    df["Open_Return_Monero_MA" + str(i_lag) + "D"] = df["Open_Return_Monero"].rolling(i_lag).mean()
    df["High_Return_Monero_MA" + str(i_lag) + "D"] = df["High_Return_Monero"].rolling(i_lag).mean()
    df["Low_Return_Monero_MA" + str(i_lag) + "D"] = df["Low_Return_Monero"].rolling(i_lag).mean()
    df["Close_Return_Monero_MA" + str(i_lag) + "D"] = df["Close_Return_Monero"].rolling(i_lag).mean()

# Ecart Type Mobile
for i_lag in [3, 6, 9]:
    df["Open_Return_MA" + str(i_lag) + "D"] = df["Open_Return"].rolling(i_lag).std()
    df["High_Return_MA" + str(i_lag) + "D"] = df["High_Return"].rolling(i_lag).std()
    df["Low_Return_MA" + str(i_lag) + "D"] = df["Low_Return"].rolling(i_lag).std()
    df["Close_Return_MA" + str(i_lag) + "D"] = df["Close_Return"].rolling(i_lag).std()

```



```

df["Open_Return_Litecoin_MA" + str(i_lag) + "D"] = df["Open_Return_Litecoin"].rolling(i_lag).skew()
df["High_Return_Litecoin_MA" + str(i_lag) + "D"] = df["High_Return_Litecoin"].rolling(i_lag).skew()
df["Low_Return_Litecoin_MA" + str(i_lag) + "D"] = df["Low_Return_Litecoin"].rolling(i_lag).skew()
df["Close_Return_Litecoin_MA" + str(i_lag) + "D"] = df["Close_Return_Litecoin"].rolling(i_lag).skew()

df["Open_Return_Monero_MA" + str(i_lag) + "D"] = df["Open_Return_Monero"].rolling(i_lag).skew()
df["High_Return_Monero_MA" + str(i_lag) + "D"] = df["High_Return_Monero"].rolling(i_lag).skew()
df["Low_Return_Monero_MA" + str(i_lag) + "D"] = df["Low_Return_Monero"].rolling(i_lag).skew()
df["Close_Return_Monero_MA" + str(i_lag) + "D"] = df["Close_Return_Monero"].rolling(i_lag).skew()

# Skew Mobile
for i_lag in [3, 6, 9]:
    df["Open_Return_MA" + str(i_lag) + "D"] = df["Open_Return"].rolling(i_lag).skew()
    df["High_Return_MA" + str(i_lag) + "D"] = df["High_Return"].rolling(i_lag).skew()
    df["Low_Return_MA" + str(i_lag) + "D"] = df["Low_Return"].rolling(i_lag).skew()
    df["Close_Return_MA" + str(i_lag) + "D"] = df["Close_Return"].rolling(i_lag).skew()

    df["Open_Return_Litecoin_MA" + str(i_lag) + "D"] = df["Open_Return_Litecoin"].rolling(i_lag).skew()
    df["High_Return_Litecoin_MA" + str(i_lag) + "D"] = df["High_Return_Litecoin"].rolling(i_lag).skew()
    df["Low_Return_Litecoin_MA" + str(i_lag) + "D"] = df["Low_Return_Litecoin"].rolling(i_lag).skew()
    df["Close_Return_Litecoin_MA" + str(i_lag) + "D"] = df["Close_Return_Litecoin"].rolling(i_lag).skew()

    df["Open_Return_Monero_MA" + str(i_lag) + "D"] = df["Open_Return_Monero"].rolling(i_lag).skew()
    df["High_Return_Monero_MA" + str(i_lag) + "D"] = df["High_Return_Monero"].rolling(i_lag).skew()
    df["Low_Return_Monero_MA" + str(i_lag) + "D"] = df["Low_Return_Monero"].rolling(i_lag).skew()
    df["Close_Return_Monero_MA" + str(i_lag) + "D"] = df["Close_Return_Monero"].rolling(i_lag).skew()

# Cosinus sur Moyenne Mobile
for i_lag in [3, 6, 9]:
    df["Open_Return_Cosinus_MA_" + str(i_lag) + "D"] = \
        df["Open_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.cos(x))
    df["High_Return_Cosinus_MA_" + str(i_lag) + "D"] = \
        df["High_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.cos(x))
    df["Low_Return_Cosinus_MA_" + str(i_lag) + "D"] = \
        df["Low_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.cos(x))
    df["Close_Return_Cosinus_MA_" + str(i_lag) + "D"] = \
        df["Close_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.cos(x))

# Sinus sur Moyenne Mobile
for i_lag in [3, 6, 9]:
    df["Open_Return_Sinus_MA_" + str(i_lag) + "D"] = \
        df["Open_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.sin(x))
    df["High_Return_Sinus_MA_" + str(i_lag) + "D"] = \
        df["High_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.sin(x))
    df["Low_Return_Sinus_MA_" + str(i_lag) + "D"] = \
        df["Low_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.sin(x))
    df["Close_Return_Sinus_MA_" + str(i_lag) + "D"] = \
        df["Close_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.sin(x))

```

```

# Exp sur Moyenne Mobile
for i_lag in [3, 6, 9]:
    df["Open_Return_Exp_MA_" + str(i_lag) + "D"] = \
        df["Open_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.exp(x))
    df["High_Return_Exp_MA_" + str(i_lag) + "D"] = \
        df["High_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.exp(x))
    df["Low_Return_Exp_MA_" + str(i_lag) + "D"] = \
        df["Low_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.exp(x))
    df["Close_Return_Exp_MA_" + str(i_lag) + "D"] = \
        df["Close_Return_MA" + str(i_lag) + "D"].apply(lambda x: np.exp(x))

# Classification
# Création de la Target : signal d'achat ou de vente
#dataset["target"] = \
#    [1 if dataset.loc[row + 1, "btc_return"] > 0 else 0
#     for row in range(0, len(dataset.loc[:, "btc_return"]))]

# Regression
# Création de la target : rentabilité du jour d'après
df["Target"] = df["Close_Return"].shift(-1)

df = df.reindex(
    np.concatenate((["Date", "Target"], df.columns[1:-1].tolist())),
    axis=1)

df = df.dropna().reset_index(drop=True)
df.head(5)

```

```

Out[124]:
      Date      Target  btc_market_price  btc_total_bitcoins  btc_market_cap \
0 2014-05-30  0.013570           609.03          12834000.0    7.816291e+09
1 2014-05-31  0.010502           620.45          12837800.0    7.965213e+09
2 2014-06-01  0.048220           674.98          12842275.0    8.668279e+09
3 2014-06-04  0.018843           644.66          12856025.0    8.287765e+09
4 2014-06-06  0.001943           655.75          12863350.0    8.435142e+09

      btc_trade_volume  btc_blocks_size  btc_avg_block_size \
0      2.175838e+07          18761.0          0.234589
1      1.080980e+07          18793.0          0.211037
2      3.394311e+07          18827.0          0.192052
3      2.831485e+07          18943.0          0.207971
4      1.286732e+07          19021.0          0.264822

      btc_n_orphaned_blocks  btc_n_transactions_per_block \
0              1.0              363.0
1              0.0              459.0
2              1.0              438.0
3              0.0              411.0

```

4		2.0		421.0
---	--	-----	--	-------

	...	Low_Return_Exp_MA_3D	Close_Return_Exp_MA_3D	\
0	...	0.288141	3.510433	
1	...	3.837640	3.049477	
2	...	3.437376	5.590915	
3	...	1.586938	0.667321	
4	...	0.182540	0.202531	

	Open_Return_Exp_MA_6D	High_Return_Exp_MA_6D	Low_Return_Exp_MA_6D	\
0	5.463303	2.221119	3.947754	
1	1.842962	4.549867	1.681664	
2	4.711021	2.076485	2.190026	
3	2.717718	2.081822	0.627698	
4	2.955675	6.960447	0.379141	

	Close_Return_Exp_MA_6D	Open_Return_Exp_MA_9D	High_Return_Exp_MA_9D	\
0	1.714983	3.271087	1.476899	
1	3.991242	1.869539	2.403294	
2	4.937080	3.260566	1.719723	
3	1.169969	3.384024	3.633603	
4	0.408644	3.902769	3.674279	

	Low_Return_Exp_MA_9D	Close_Return_Exp_MA_9D
0	2.654409	1.773699
1	1.579691	3.007898
2	2.268653	2.886790
3	1.148574	2.199385
4	0.667117	1.764899

[5 rows x 127 columns]

- Normalisation des données

```
In [125]: dataset = df
          # Min-Max Scalings
          dataset.iloc[:, 1:] = dataset.iloc[:, 1:].apply(lambda x: (x - np.min(x))/(np.max(x) - np.min(x)))

          dataset.head(10)
```

```
Out[125]:
```

	Date	Target	btc_market_price	btc_total_bitcoins	btc_market_cap	\
0	2014-05-30	0.485048	0.022385	0.000000	0.016643	
1	2014-05-31	0.478435	0.022976	0.000940	0.017103	
2	2014-06-01	0.559737	0.025798	0.002048	0.019272	
3	2014-06-04	0.496414	0.024229	0.005450	0.018098	
4	2014-06-06	0.459986	0.024803	0.007263	0.018552	
5	2014-06-07	0.459648	0.024728	0.008209	0.018503	
6	2014-06-08	0.432868	0.024609	0.009137	0.018419	

7	2014-06-09	0.469047	0.024554	0.009979	0.018383
8	2014-06-10	0.389366	0.024500	0.011074	0.018351
9	2014-06-11	0.298925	0.023987	0.011977	0.017965

	btc_trade_volume	btc_blocks_size	btc_avg_block_size	\
0	0.003575	0.000000	0.073250	
1	0.001528	0.000231	0.048326	
2	0.005853	0.000476	0.028235	
3	0.004801	0.001312	0.045082	
4	0.001913	0.001874	0.105245	
5	0.001767	0.002090	0.037778	
6	0.000632	0.002299	0.034673	
7	0.002094	0.002544	0.094944	
8	0.001982	0.002832	0.066179	
9	0.002705	0.003092	0.091972	

	btc_n_orphaned_blocks	btc_n_transactions_per_block	\
0	0.142857	0.041436	
1	0.000000	0.080435	
2	0.142857	0.071904	
3	0.000000	0.060935	
4	0.285714	0.064998	
5	0.285714	0.055248	
6	0.285714	0.051186	
7	0.142857	0.030061	
8	0.285714	0.011375	
9	0.142857	0.024374	

	...	Low_Return_Exp_MA_3D	Close_Return_Exp_MA_3D	\
0	...	0.020313	0.608826	
1	...	0.668587	0.524638	
2	...	0.595484	0.988801	
3	...	0.257523	0.089566	
4	...	0.001026	0.004677	
5	...	0.408491	0.027083	
6	...	0.236051	0.999488	
7	...	0.942956	0.000030	
8	...	0.051654	0.019669	
9	...	0.486011	0.106510	

	Open_Return_Exp_MA_6D	High_Return_Exp_MA_6D	Low_Return_Exp_MA_6D	\
0	0.500525	0.191485	0.368380	
1	0.162842	0.400767	0.151991	
2	0.430357	0.178487	0.200535	
3	0.244434	0.178967	0.051348	
4	0.266629	0.617403	0.027614	
5	0.318354	0.741408	0.060061	
6	0.445188	0.339706	0.077338	

7	0.568597	0.247656	0.095129
8	0.141806	0.091603	0.130587
9	0.089555	0.059537	0.306438

	Close_Return_Exp_MA_6D	Open_Return_Exp_MA_9D	High_Return_Exp_MA_9D	\
0	0.152270	0.219382	0.088035	
1	0.365816	0.123296	0.145602	
2	0.454549	0.218661	0.103124	
3	0.101140	0.227125	0.222055	
4	0.029717	0.262688	0.224583	
5	0.043875	0.297693	0.244211	
6	0.057650	0.308350	0.239730	
7	0.017751	0.346795	0.722995	
8	0.017846	0.389390	0.809245	
9	0.030135	0.494066	0.167831	

	Low_Return_Exp_MA_9D	Close_Return_Exp_MA_9D
0	0.164014	0.118860
1	0.096204	0.204594
2	0.139674	0.196181
3	0.069002	0.148430
4	0.038625	0.118249
5	0.053619	0.128334
6	0.043217	0.124655
7	0.055810	0.049512
8	0.062254	0.056779
9	0.104012	0.077423

[10 rows x 127 columns]

6.2 Estimation du modèle

- Séparation du dataset en train / test set

In [126]: *# Splitting X in training and testing set*

```
train_set, test_set = skl_md1_selection.train_test_split(dataset,
                                                         train_size=0.66,
                                                         test_size=0.34)
```

```
X_train = train_set.iloc[:, 2:]
Y_train = train_set.iloc[:, 1]
X_test = test_set.iloc[:, 2:]
Y_test = test_set.iloc[:, 1]
```

```
X_train.head(5)
```

Out[126]:

	btc_market_price	btc_total_bitcoins	btc_market_cap	btc_trade_volume	\
417	0.004602	0.406774	0.004383	0.002842	

625	0.011619	0.605905	0.011435	0.009023
31	0.024197	0.035176	0.018312	0.001447
949	0.037084	0.814424	0.036931	0.005050
595	0.010794	0.577429	0.010546	0.007294

	btc_blocks_size	btc_avg_block_size	btc_n_orphaned_blocks	\
417	0.164505	0.218889	0.428571	
625	0.302944	0.643204	0.000000	
31	0.008806	0.124665	0.142857	
949	0.582884	0.870451	0.142857	
595	0.278031	0.494704	0.000000	

	btc_n_transactions_per_block	btc_median_confirmation_time	\
417	0.220992	0.105634	
625	0.379830	0.116197	
31	0.069466	0.109155	
949	0.722588	0.295070	
595	0.448484	0.054930	

	btc_hash_rate	...	Low_Return_Exp_MA_3D	\
417	0.010771	...	0.002852	
625	0.047577	...	0.031828	
31	0.001262	...	0.928016	
949	0.111808	...	0.110487	
595	0.040191	...	0.031696	

	Close_Return_Exp_MA_3D	Open_Return_Exp_MA_6D	High_Return_Exp_MA_6D	\
417	0.004956	0.005697	0.002987	
625	0.999579	0.062438	0.065590	
31	0.255520	0.335948	0.564966	
949	0.042182	0.146475	0.208309	
595	0.653416	0.036335	0.937034	

	Low_Return_Exp_MA_6D	Close_Return_Exp_MA_6D	Open_Return_Exp_MA_9D	\
417	0.002543	0.009349	0.004433	
625	0.104093	0.167555	0.043817	
31	0.527932	0.576896	0.118522	
949	0.025065	0.068369	0.417839	
595	0.104486	0.324257	0.015288	

	High_Return_Exp_MA_9D	Low_Return_Exp_MA_9D	Close_Return_Exp_MA_9D
417	0.005271	0.003627	0.005421
625	0.020711	0.041174	0.061209
31	0.231817	0.159569	0.118191
949	0.433807	0.029055	0.334613
595	0.372724	0.062221	0.161905

[5 rows x 125 columns]

A) On estime la Régression Elastic Net pour alpha variant de 0 à 1

```
In [127]: # ATTENTION : dans la fonction ElasticNet de sklearn, le paramètre
# l1_ratio correspond au alpha de la formule du projet,
# tandis que le paramètre alpha correspond au lambda de la formule du projet ...

all_beta_chap = dict()
#models = list()
alpha_range = [0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1]
for i_alpha in alpha_range:
    en = skl_linear_md1.ElasticNet(l1_ratio=i_alpha, fit_intercept=False, max_iter =
    en.fit(X_train, Y_train)
    beta_chap = en.coef_
    all_beta_chap.update({i_alpha: beta_chap})
    #models.append(en)

beta_vs_alpha = pd.DataFrame(all_beta_chap)


plt.figure()
plt.suptitle("Regularization Path for Elastic Net over Alpha Parameter",
            fontsize=14)
plt.title("Lambda = 1", fontsize=10)
for row in beta_vs_alpha.iterrows():
    row = row[1]
    plt.plot(alpha_range, row)
plt.ylabel("Coefficients values")
plt.xlabel("Alpha")

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_desc
ConvergenceWarning)

Out[127]: Text(0.5,0,'Alpha')
```

Regularization Path for Elastic Net over Alpha Parameter

Lambda = 1




B) Lasso Regularization Path

```
In [142]: print_regularization_path(X_train,Y_train,isLasso=True)
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:14: UserWarning
```


```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:141: UserWarning:
  positive)
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:141: ConvergenceWarning)
```

```
In [143]: lasso = skl_linear_mdl.ElasticNet(l1_ratio=1, fit_intercept= False, max_iter = 10000)
lasso_path = lasso.path(X_train, Y_train, l1_ratio=1)
lasso_path_result = lasso_path[1]
plt.figure()
plt.title("Lasso Regularization Path")
for i in range(0, len(lasso_path_result)):
    plt.plot(lasso_path[0], lasso_path_result[i])
plt.ylabel("Coefficients values")
plt.xlabel("Lambda")
```

```
Out[143]: Text(0.5,0,'Lambda')
```



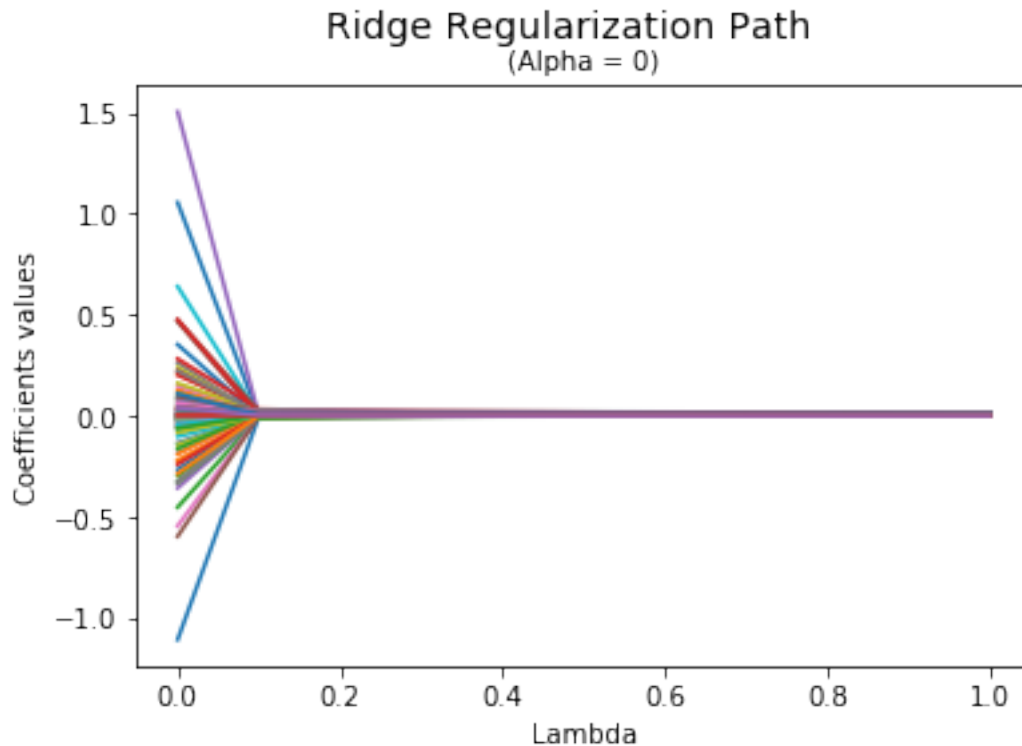
B.1) Ridge Regularization Path

```
In [144]: print_regularization_path(X_train,Y_train,isLasso=False)
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:14: UserWarning
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
    positive)
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
    ConvergenceWarning)
```



```
In [145]: # Méthode 1.1) Avec une boucle
all_beta_chap = dict()
lambda_range = [0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1]
for i_lambda in lambda_range:
    ridge = skl_linear_mdl.ElasticNet(alpha=i_lambda, l1_ratio=0, fit_intercept=False)
    ridge.fit(X_train, Y_train)
    beta_chap = ridge.coef_
    all_beta_chap.update({i_lambda: beta_chap})

beta_vs_lambda = pd.DataFrame(all_beta_chap)

# Méthode 1.2) Avec la méthode Path
ridge = skl_linear_mdl.ElasticNet(l1_ratio=1, fit_intercept=False, max_iter = 10000)
lasso_path = ridge.path(X_train, Y_train, l1_ratio=1)


plt.figure()
plt.suptitle("Ridge Regularization Path", fontsize=14)
plt.title("(Alpha = 0)", fontsize=10)
for row in beta_vs_lambda.iterrows():
    row = row[1]
    plt.plot(lambda_range, row)
plt.ylabel("Coefficients values")
plt.xlabel("Lambda")
```

```

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:6: UserWarning
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)

```

```
Out[145]: Text(0.5,0,'Lambda')
```



Optimisation du / des paramètres de régularisation (lambda et alpha)

- Lasso

```

In [159]: lambda_opt_lasso, _, all_quad_risk = optimise_params(X_train, Y_train, X_test, Y_test)
          print_plot("Lasso Regression", "Lambda Optimal : " + str(lambda_opt_lasso), "Lambda", "Lasso")
          plt.plot(all_quad_risk.columns,
                   all_quad_risk.values.reshape(11, 1))


```

```

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning
if __name__ == '__main__':
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
ConvergenceWarning)

```

Out[159]: [<matplotlib.lines.Line2D at 0x1a12ca7470>]




- Ridge

```
In [160]: lambda_opt_ridge, _, all_quad_risk = optimise_params(X_train, Y_train, X_test, Y_test)
print_plot("Ridge Regression", "Lambda Optimal : " + str(lambda_opt_ridge), "Lambda", "MSE")
plt.plot(all_quad_risk.columns,
         all_quad_risk.values.reshape(11, 1))
```

```
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning:
  if __name__ == '__main__':
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:111: UserWarning:
  positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:111: UserWarning:
  ConvergenceWarning)
```

Out[160]: [<matplotlib.lines.Line2D at 0x1a12b8c9e8>]



- Elastic Net


```
In [161]: lambda_opt_elastic_net, alpha_opt_elastic_net, all_quad_risk = optimise_params(X_train, y_train,
print_plot3D("Elastic Net Regression\nAlpha Optimal : " + str(alpha_opt_elastic_net)
            "\nLambda Optimal : " + str(lambda_opt_elastic_net), "", "Alpha", "Lambda",
            np.linspace(0, 1, 11), np.linspace(0, 1, 11), all_quad_risk, 0)

print_plot3D("Elastic Net Regression\nAlpha Optimal : " + str(alpha_opt_elastic_net)
            "\nLambda Optimal : " + str(lambda_opt_elastic_net), "", "Lambda", "Alpha",
            np.linspace(0, 1, 11), np.linspace(0, 1, 11), all_quad_risk, 1)


opt_elastic_net=[]
opt_elastic_net.append(alpha_opt_elastic_net)
opt_elastic_net.append(lambda_opt_elastic_net)

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:9: UserWarning
  if __name__ == '__main__':
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100:
  positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100:
  ConvergenceWarning)
```

Elastic Net Regression
Alpha Optimal : 0.0
Lambda Optimal : 0.1



Elastic Net Regression
 Alpha Optimal : 0.0
 Lambda Optimal : 0.1



- Score des estimateurs

```
In [162]: print_score(X_train,Y_train,lambda_opt_lasso,lambda_opt_ridge,opt_elastic_net)

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:8: UserWarning

/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  positive)
/Users/virgileamato/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:
  ConvergenceWarning)
```


	Score (R2)	MSE
Lasso	-0.356518	0.00899991
Ridge	-0.343664	0.00891463
Elastic Net	-0.343664	0.00891463
BEST MODEL	Ridge	Ridge